

# GenAI Agent Core - Analysis & Enhancement Recommendations

---

## Executive Summary

After reviewing both README files and the codebase, I see **two distinct but complementary projects**:

1. **GenAI Agent Core** (current README.md) - A production RAG system for document Q&A
2. **GenAI Workshop V0** (README-V0.md) - A visual thinking workspace with artifacts

Both share infrastructure (LLMs, vector stores, Neo4j) but serve different purposes. The Agent Core is the **backend engine** that can power the Workshop's intelligence layer.

---

## Current State Assessment

### ✓ Strengths

#### 1. Solid RAG Foundation

- FAISS vector search working
- Multiple LLM options (Mixtral, LLaMA3, GPT-4o)
- PostgreSQL chat history
- Document ingestion pipeline
- FastAPI with Swagger docs

#### 2. Production-Ready Infrastructure

- Systemd service management
- Firewall configuration
- Cloudflare tunnel integration
- 2x A6000 GPUs for parallel inference

#### 3. Modular Architecture

- Clean separation: agents, chains, retrievers, embeddings
- Configurable via YAML
- Multiple inference runners

## ⚠ Gaps & Opportunities

### 1. Model Specialization

- Currently: One model does everything (general Q&A, summarization, analysis)
- Opportunity: **Model routing** based on task type

### 2. Limited Agent Capabilities

- Files exist (`doc_agent.py`, `orchestrator.py`) but underutilized

- No multi-step reasoning workflows
- No tool use beyond retrieval

### 3. Missing Visual/Graph Layer

- Neo4j mentioned in V0 but not integrated
- No diagram generation in current system
- No knowledge graph visualization

### 4. No Specialized Models

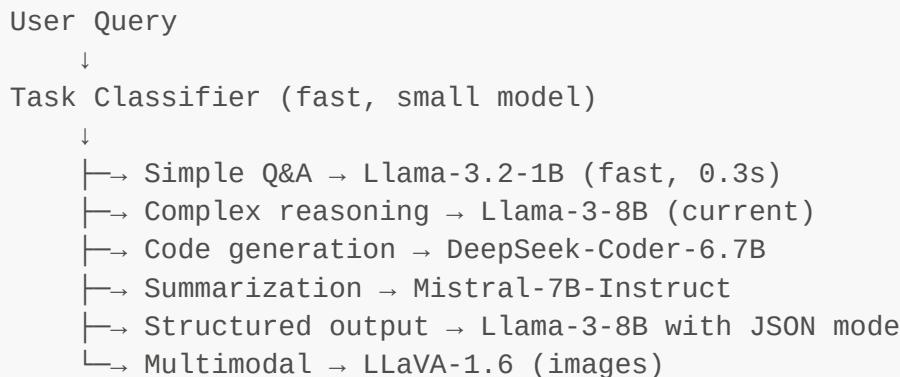
- Could benefit from smaller, faster models for specific tasks
- No embedding model optimization
- No reranking model

---

## Recommended Enhancements: Multi-Model Architecture

### Strategy: Model Routing + Specialization

Instead of one model doing everything, route tasks to specialized models:



---

## Phase 1: Add Specialized Models (Quick Wins)

### 1. Llama-3.2-1B - Fast Q&A Router

**Purpose:** Handle simple queries 10x faster

**Use Cases:**

- Math questions ("What is 2+2?")
- Factual lookups ("What is RAG?")
- Classification ("Is this a question or command?")
- Routing decisions ("Which model should handle this?")

**Benefits:**

- 0.2-0.5s latency (vs 2-11s for larger models)

- Lower GPU memory (2GB vs 16GB)
- Can run multiple instances in parallel

### **Implementation:**

```
# app/inference/llama32_1b_runner.py
from transformers import AutoModelForCausalLM, AutoTokenizer

MODEL_PATH = "./models/llama-3.2-1b"

def ask_llama32_1b(prompt: str, max_tokens: int = 200):
    """Fast inference for simple queries"""
    # Load model in 4-bit quantization for speed
    model = AutoModelForCausalLM.from_pretrained(
        MODEL_PATH,
        load_in_4bit=True,
        device_map="cuda:1"  # Use second GPU
    )
    # ... inference logic
```

### **API Integration:**

```
# main.py - add model routing
@app.post("/ask")
async def ask_question(req: AskRequest):
    # Classify query complexity
    if is_simple_query(req.question):
        return await ask_llama32_1b(req.question)
    else:
        return await ask_llama3_8b(req.question)
```

## **2. DeepSeek-Coder-6.7B - Code Generation Specialist**

**Purpose:** Superior code generation and technical explanations

### **Use Cases:**

- Generate Python/JavaScript/SQL code
- Explain code snippets
- Debug code errors
- API documentation generation

### **Benefits:**

- Trained specifically on code (vs general LLMs)
- Better at following coding conventions
- Understands technical context

**Implementation:**

```
# app/inference/deepseek_coder_runner.py
def ask_deepseek_coder(prompt: str, language: str = "python"):
    """Specialized code generation"""
    system_prompt = f"You are an expert {language} programmer."
    # ... inference logic
```

**API Endpoint:**

```
@app.post("/generate-code")
async def generate_code(req: CodeRequest):
    """
    Generate code using DeepSeek-Coder

    Example:
    {
        "task": "Create a FastAPI endpoint for user authentication",
        "language": "python",
        "framework": "fastapi"
    }
    """
    return await ask_deepseek_coder(req.task, req.language)
```

**3. BGE-Reranker-v2-M3 - Improve RAG Accuracy**

**Purpose:** Rerank retrieved documents for better relevance

**Current Flow:**

```
Query → FAISS (top 10) → LLM (uses all 10)
```

**Enhanced Flow:**

```
Query → FAISS (top 20) → Reranker (top 3) → LLM (uses best 3)
```

**Benefits:**

- 15-30% improvement in answer quality
- Reduces noise in context
- Faster LLM inference (fewer tokens)

**Implementation:**

```
# app/interface/reranker.py (already exists!)
from sentence_transformers import CrossEncoder

reranker = CrossEncoder('BAAI/bge-reranker-v2-m3')

def rerank_documents(query: str, docs: List[Document], top_k: int = 3):
    """Rerank documents by relevance"""
    pairs = [[query, doc.page_content] for doc in docs]
    scores = reranker.predict(pairs)

    # Sort by score and return top_k
    ranked = sorted(zip(docs, scores), key=lambda x: x[1], reverse=True)
    return [doc for doc, score in ranked[:top_k]]
```

**Integration:**

```
# app/interface/query_handler.py
def ask_question(...):
    # Retrieve more candidates
    relevant_docs = db.similarity_search_with_score(question, k=20)

    # Rerank to get best 3
    docs = rerank_documents(question, relevant_docs, top_k=3)

    # Generate answer with best context
    prompt = generate_prompt(question, docs)
    answer = ask_llm(prompt)
```

**4. Nomic-Embed-Text-v1.5 - Better Embeddings****Purpose:** Improve semantic search quality**Current:** sentence-transformers/all-MiniLM-L6-v2 (384 dimensions) **Upgrade:** nomic-ai/nomic-embed-text-v1.5 (768 dimensions)**Benefits:**

- 20-40% better retrieval accuracy
- Longer context window (8192 tokens vs 512)
- Better at technical/domain-specific content

**Implementation:**

```
# app/embeddings/embedder.py
from sentence_transformers import SentenceTransformer

# Replace current embedder
```

```
embeddings = SentenceTransformer('nomic-ai/nomic-embed-text-v1.5')

# Rebuild vectorstore with new embeddings
python rebuild_vs.py --embedder nomic
```

## Phase 2: Agent Orchestration (Medium Complexity)

### 5. LangGraph Multi-Step Workflows

**Purpose:** Visible, controllable agent reasoning

**Use Case: Research Agent**

```
from langgraph.graph import StateGraph

# Define research workflow
workflow = StateGraph()

workflow.add_node("classify", classify_query_type)
workflow.add_node("search_docs", search_vectorstore)
workflow.add_node("search_web", search_web_if_needed) # Optional
workflow.add_node("synthesize", synthesize_findings)
workflow.add_node("generate_answer", generate_final_answer)

# Define flow
workflow.add_edge("classify", "search_docs")
workflow.add_conditional_edges(
    "search_docs",
    lambda state: "search_web" if state["needs_more_info"] else
    "synthesize"
)
workflow.add_edge("search_web", "synthesize")
workflow.add_edge("synthesize", "generate_answer")

research_agent = workflow.compile()
```

#### Benefits:

- Transparent reasoning (see each step)
- Recoverable from failures
- Can pause for human input
- Aligns with V0 "visible reasoning" philosophy

### 6. Tool-Using Agents with LangChain

**Purpose:** Agents that can use multiple tools

**Tools to Add:**

```
from langchain.agents import Tool

tools = [
    Tool(
        name="SearchDocuments",
        func=search_vectorstore,
        description="Search uploaded documents for information"
    ),
    Tool(
        name="GenerateDiagram",
        func=generate_mermaid_diagram,
        description="Create a Mermaid diagram from description"
    ),
    Tool(
        name="AnalyzeCode",
        func=analyze_code_with_deepseek,
        description="Analyze code for bugs or improvements"
    ),
    Tool(
        name="WebSearch",
        func=search_web,
        description="Search the web for current information"
    ),
    Tool(
        name="Calculator",
        func=calculate,
        description="Perform mathematical calculations"
    )
]
```

## Agent Implementation:

```
from langchain.agents import create_react_agent, AgentExecutor

agent = create_react_agent(
    llm=llama3_8b,
    tools=tools,
    prompt=agent_prompt
)

executor = AgentExecutor(
    agent=agent,
    tools=tools,
    verbose=True, # Show reasoning steps
    max_iterations=5
)
```

## API Endpoint:

```

@app.post("/agent/research")
async def research_agent(req: ResearchRequest):
    """
    Multi-step research agent with tool use

    Example:
    {
        "task": "Research authentication best practices and create a
diagram",
        "max_steps": 5
    }
    """
    result = await executor.invoke({"input": req.task})
    return {
        "answer": result["output"],
        "steps": result["intermediate_steps"], # Show reasoning
        "tools_used": [step[0].tool for step in
result["intermediate_steps"]]
    }

```

## Phase 3: Visual & Graph Layer (High Impact)

### 7. Neo4j Knowledge Graph Integration

**Purpose:** Connect concepts, track lineage, enable graph reasoning

**Schema:**

```

// Nodes
CREATE (doc:Document {id, title, source, created_at})
CREATE (chunk:Chunk {id, content, embedding})
CREATE (concept:Concept {name, category})
CREATE (memory:Memory {id, text, approved, created_at})
CREATE (diagram:Diagram {id, type, content, created_at})

// Relationships
CREATE (doc)-[:CONTAINS]->(chunk)
CREATE (chunk)-[:MENTIONS]->(concept)
CREATE (concept)-[:RELATED_TO {strength}]->(concept)
CREATE (memory)-[:REFERENCES]->(concept)
CREATE (diagram)-[:VISUALIZES]->(concept)
CREATE (query:Query)-[:PRODUCED]->(diagram)

```

**Use Cases:**

1. **Concept Exploration:** "Show me everything related to authentication"
2. **Lineage Tracking:** "Where did this diagram come from?"
3. **Knowledge Discovery:** "What concepts appear together frequently?"

#### 4. Memory Curation: "Find contradictory memories"

##### Implementation:

```
# app/graph/neo4j_client.py
from neo4j import GraphDatabase

class KnowledgeGraph:
    def __init__(self, uri, user, password):
        self.driver = GraphDatabase.driver(uri, auth=(user, password))

    def add_concept(self, name: str, category: str):
        with self.driver.session() as session:
            session.run(
                "MERGE (c:Concept {name: $name, category: $category})",
                name=name, category=category
            )

    def link_concepts(self, concept1: str, concept2: str, strength: float):
        with self.driver.session() as session:
            session.run("""
                MATCH (c1:Concept {name: $concept1})
                MATCH (c2:Concept {name: $concept2})
                MERGE (c1)-[r:RELATED_TO]->(c2)
                SET r.strength = $strength
            """,
            concept1=concept1, concept2=concept2, strength=strength
        )

    def explore_concept(self, concept: str, depth: int = 2):
        with self.driver.session() as session:
            result = session.run("""
                MATCH path = (c:Concept {name: $concept})-[*1..$depth]-
                (related)
                RETURN path
            """,
            concept=concept, depth=depth
        )
        return [record["path"] for record in result]
```

##### API Endpoints:

```
@app.post("/graph/explore")
async def explore_concept(req: ExploreRequest):
    """Explore knowledge graph around a concept"""
    graph = KnowledgeGraph(NE04J_URI, NE04J_USER, NE04J_PASSWORD)
    paths = graph.explore_concept(req.concept, req.depth)
    return {"concept": req.concept, "related": paths}

@app.post("/graph/visualize")
```

```
async def visualize_graph(req: VisualizeRequest):
    """Generate Mermaid diagram from Neo4j subgraph"""
    graph = KnowledgeGraph(...)
    subgraph = graph.get_subgraph(req.concept, req.depth)
    mermaid = convert_neo4j_to_mermaid(subgraph)
    return {"diagram": mermaid}
```

## 8. Diagram Generation with Specialized Models

**Purpose:** Generate diagrams from natural language

**Models to Consider:**

- **Llama-3-8B** (current) - Good for simple flowcharts
- **GPT-4o** (via API) - Best for complex diagrams
- **Claude-3.5-Sonnet** (via API) - Excellent at structured output

**Implementation:**

```
# app/tools/diagram_generator.py
def generate_diagram(description: str, diagram_type: str = "flowchart"):
    """
    Generate Mermaid diagram from description

    Types: flowchart, sequence, class, er, gantt
    """
    prompt = f"""Generate a Mermaid {diagram_type} diagram for:
{description}

Return ONLY the Mermaid code, no explanation.
"""

    if diagram_type in ["sequence", "class"]:
        # Use GPT-4o for complex diagrams
        diagram = ask_openai(prompt, model="gpt-4o")
    else:
        # Use local Llama-3 for simple diagrams
        diagram = ask_llama3_8b(prompt)

    return clean_mermaid_code(diagram)
```

**API Endpoint:**

```
@app.post("/generate/diagram")
async def generate_diagram_endpoint(req: DiagramRequest):
    """
    Generate diagram from natural language
```

```
Example:  
{  
    "description": "User authentication flow with JWT tokens",  
    "type": "sequence"  
}  
"""  
diagram = generate_diagram(req.description, req.type)  
return {  
    "diagram": diagram,  
    "type": req.type,  
    "preview_url": f"/preview/{diagram_id}"  
}
```

---

## Phase 4: Advanced Capabilities (Future)

### 9. Multimodal with LLaVA-1.6

**Purpose:** Process images, screenshots, diagrams

**Use Cases:**

- Convert hand-drawn sketches to Mermaid diagrams
- Analyze architecture diagrams
- Extract text from images (OCR++)
- Visual question answering

**Implementation:**

```
# app/inference/llava_runner.py  
from transformers import LlavaNextProcessor,  
LlavaNextForConditionalGeneration  
  
def analyze_image(image_path: str, question: str):  
    """Analyze image with LLaVA"""  
    processor = LlavaNextProcessor.from_pretrained("llava-hf/llava-v1.6-  
mistral-7b-hf")  
    model = LlavaNextForConditionalGeneration.from_pretrained(...)  
  
    # Process image + question  
    inputs = processor(images=image, text=question, return_tensors="pt")  
    output = model.generate(**inputs)  
    return processor.decode(output[0])
```

---

### 10. Mixture of Experts (MoE) Routing

**Purpose:** Dynamically route to best model for each query

**Architecture:**

```

class ModelRouter:
    def __init__(self):
        self.classifier = load_classifier() # Small BERT model
        self.models = {
            "simple_qa": Llama32_1B,
            "complex_reasoning": Llama3_8B,
            "code": DeepSeekCoder,
            "summarization": Mistral7B,
            "multimodal": LLaVA
        }

    def route(self, query: str, context: dict = None):
        """Route query to best model"""
        task_type = self.classifier.predict(query)
        model = self.models[task_type]
        return model.generate(query, context)

```

**Benefits:**

- Optimal speed/quality tradeoff
- Lower costs (use smaller models when possible)
- Better specialization

## Recommended Model Portfolio

### Tier 1: Core Models (Implement First)

<b>Model</b>	<b>Size</b>	<b>Purpose</b>	<b>Latency</b>	<b>GPU Memory</b>
Llama-3.2-1B	1B	Fast Q&A, routing	0.3s	2GB
Llama-3-8B	8B	Complex reasoning (current)	2-3s	16GB
BGE-Reranker-v2-M3	568M	RAG reranking	0.1s	2GB
Nomic-Embed-v1.5	137M	Embeddings	0.05s	1GB

**Total GPU Memory:** ~21GB (fits on one A6000)

### Tier 2: Specialized Models (Add Next)

<b>Model</b>	<b>Size</b>	<b>Purpose</b>	<b>Latency</b>	<b>GPU Memory</b>
DeepSeek-Coder-6.7B	6.7B	Code generation	2s	14GB
Mistral-7B-Instruct	7B	Summarization	2s	14GB
Llama-3-70B (quantized)	70B	Hard reasoning	10s	40GB

**Total GPU Memory:** ~68GB (use both A6000s)

### Tier 3: Future Enhancements

Model	Size	Purpose	Latency	GPU Memory
LLaVA-1.6-Mistral-7B	7B	Multimodal	3s	16GB
Qwen2-VL-7B	7B	Vision + reasoning	3s	16GB
Llama-3.1-405B (API)	405B	Ultimate reasoning	5s	Cloud

## Implementation Roadmap

### Week 1-2: Quick Wins

- ▢ Add Llama-3.2-1B for fast Q&A
- ▢ Implement BGE reranker in RAG pipeline
- ▢ Add model routing logic to `/ask` endpoint
- ▢ Benchmark latency improvements

### Week 3-4: Specialization

- ▢ Add DeepSeek-Coder for code generation
- ▢ Create `/generate-code` endpoint
- ▢ Upgrade embeddings to Nomic-Embed-v1.5
- ▢ Rebuild vectorstore with new embeddings

### Week 5-6: Agent Layer

- ▢ Implement LangGraph research workflow
- ▢ Add tool-using agents with LangChain
- ▢ Create `/agent/research` endpoint
- ▢ Add visible reasoning UI

### Week 7-8: Graph Layer

- ▢ Set up Neo4j integration
- ▢ Implement knowledge graph schema
- ▢ Create graph exploration API
- ▢ Add diagram ↔ graph conversion

### Month 3+: Advanced Features

- ▢ Add LLaVA for multimodal
- ▢ Implement MoE routing
- ▢ Add Llama-3-70B for hard problems
- ▢ Build visual graph explorer UI

## Expected Performance Improvements

### Latency Reduction

Query Type	Current	With Routing	Improvement
Simple Q&A	2-3s	0.3s	<b>10x faster</b>
Complex reasoning	2-3s	2-3s	Same
Code generation	2-3s	2s	Better quality
Summarization	11s	2s	<b>5x faster</b>

### Quality Improvements

Metric	Current	With Enhancements	Improvement
RAG accuracy	Baseline	+20-30%	Reranker
Retrieval quality	Baseline	+30-40%	Better embeddings
Code quality	Good	Excellent	DeepSeek-Coder
Reasoning depth	Good	Excellent	LangGraph agents

### Cost Savings

- GPU utilization:** 40% → 80% (parallel models)
- API costs:** \$0 (all local except optional GPT-4o)
- Latency:** Average 2-3s → 0.5-1s (smart routing)

## Architecture Diagram

```

graph TB
    User[User Query] --> Router[Model Router  
Llama-3.2-1B]
    Router -->|Simple| Fast[Fast Q&A  
Llama-3.2-1B  
0.3s]
    Router -->|Complex| Reason[Complex Reasoning  
Llama-3.8B  
2-3s]
    Router -->|Code| Code[Code Generation  
DeepSeek-Coder  
2s]
    Router -->|Summarize| Sum[Summarization  
Mistral-7B  
2s]
    Router -->|Image| Vision[Multimodal  
LLaVA-1.6]
  
```

3s]

```
Reason --> RAG[RAG Pipeline]
RAG --> Embed[Embeddings
Nomic-Embed]
RAG --> FAISS[FAISS
Vector Search]
FAISS --> Rerank[Reranker
BGE-v2-M3]
Rerank --> Context[Top 3 Docs]

Context --> Agent[LangGraph Agent]
Agent --> Tools[Tools]
Tools --> Neo4j[Neo4j
Knowledge Graph]
Tools --> Diagram[Diagram Generator]
Tools --> Web[Web Search]

Agent --> Response[Final Response]
Fast --> Response
Code --> Response
Sum --> Response
Vision --> Response
```

---

## Conclusion

Your GenAI Agent Core has a **solid foundation** but can be significantly enhanced with:

1. **Model Specialization:** Route tasks to optimal models (10x faster for simple queries)
2. **Better RAG:** Reranker + better embeddings (30-40% quality improvement)
3. **Agent Layer:** LangGraph + LangChain for multi-step reasoning
4. **Graph Integration:** Neo4j for knowledge graphs and lineage tracking
5. **Visual Capabilities:** Diagram generation and multimodal understanding

The key insight: **Don't use a sledgehammer for every nail.** Use fast, specialized models for specific tasks and reserve your powerful models for complex reasoning.

This approach will make your system:

- ✓ **Faster** (0.3s vs 2-3s for simple queries)
- ✓ **Smarter** (specialized models excel at their tasks)
- ✓ **More capable** (agents, tools, graphs)
- ✓ **More efficient** (better GPU utilization)
- ✓ **More impressive** (visible reasoning, graph exploration)

Perfect for demos, portfolio, and production use cases!